

ST 440 Lab #1

Christian Nielsen

(1) Chapter 1, Problem 4

X_1 and X_2 have joint PMF

x_1	x_2	$\text{Prob}(X_1 = x_1, X_2 = x_2)$
0	0	0.15
1	0	0.15
2	0	0.15
0	1	0.15
1	1	0.20
2	1	0.20

(a) Compute the marginal distribution of X_1 .

The marginal distribution is computed by summing over the other variable in the joint PMF.

x_1	$\text{Prob}(X_1 = x_1)$
0	0.3
1	0.35
2	0.35

(b) Compute the marginal distribution of X_2 .

x_2	$\text{Prob}(X_2 = x_2)$
0	0.45
1	0.55

(c) Compute the conditional distribution of $X_1 \mid X_2$.

The general expression for the conditional distributions of $X_1 \mid X_2 = x_2$ is

$$f_{1|2}(x_1 \mid X_2 = x_2) = \frac{f(x_1, x_2)}{f_2(x_2)}.$$

So the conditional distribution of $X_1 \mid X_2$ in this case is

$$\begin{aligned}
f_{1|2}(0 | X_2 = 0) &= \frac{f(0,0)}{f_2(0)} = \frac{0.15}{0.45} = 0.3333, \\
f_{1|2}(1 | X_2 = 0) &= \frac{f(1,0)}{f_2(0)} = \frac{0.15}{0.45} = 0.3333, \\
f_{1|2}(2 | X_2 = 0) &= \frac{f(2,0)}{f_2(0)} = \frac{0.15}{0.45} = 0.3333, \\
f_{1|2}(0 | X_2 = 1) &= \frac{f(0,1)}{f_2(1)} = \frac{0.15}{0.55} = 0.2727, \\
f_{1|2}(1 | X_2 = 1) &= \frac{f(1,1)}{f_2(1)} = \frac{0.2}{0.55} = .3636, \\
f_{1|2}(2 | X_2 = 1) &= \frac{f(2,1)}{f_2(1)} = \frac{0.2}{0.55} = 0.3636
\end{aligned}$$

Under each condition ($X_2 = 0$ and $X_2 = 1$), the conditional probabilities sum to one, which is required by definition.

(d) Compute the conditional distribution of $X_2 | X_1$.

The general expression for the conditional distributions of $X_2 | X_1 = x_1$ is

$$f_{2|1}(x_2 | X_1 = x_1) = \frac{f(x_1, x_2)}{f_1(x_1)}.$$

So the conditional distribution of $X_2 | X_1$ in this case is

$$\begin{aligned}
f_{2|1}(0 | X_1 = 0) &= \frac{f(0,0)}{f_1(0)} = \frac{0.15}{0.3} = 0.5, \\
f_{2|1}(1 | X_1 = 0) &= \frac{f(0,1)}{f_1(0)} = \frac{0.15}{0.3} = 0.5, \\
f_{2|1}(0 | X_1 = 1) &= \frac{f(1,0)}{f_1(1)} = \frac{0.15}{0.35} = 0.4286, \\
f_{2|1}(1 | X_1 = 1) &= \frac{f(1,1)}{f_1(1)} = \frac{0.2}{0.35} = 0.5714, \\
f_{2|1}(0 | X_1 = 2) &= \frac{f(2,0)}{f_1(2)} = \frac{0.15}{0.35} = 0.4286, \\
f_{2|1}(1 | X_1 = 2) &= \frac{f(2,1)}{f_1(2)} = \frac{0.2}{0.35} = 0.5714,
\end{aligned}$$

Under each condition ($X_1 = 0, X_1 = 1$, and $X_1 = 2$), the conditional probabilities sum to one, which is required by definition.

(e) Are X_1 and X_2 independent? Justify your answer.

X_1 and X_2 are independent if and only if

$$f(x_1, x_2) = f_1(x_1)f_2(x_2).$$

We can check if this is true using an example from the tables above:

$$f(0, 0) = 0.15 \neq 0.135 = (0.3)(0.45) = f_1(0)f_2(0)$$

This violates the definition for independence, so X_1 and X_2 are **dependent**.

(2) Chapter 1, Problem 6

Assume (X_1, X_2) have bivariate PDF

$$f(x_1, x_2) = \frac{1}{2\pi}(1 + x_1^2 + x_2^2)^{-3/2}$$

- (a) Plot the conditional distribution of $X_1 \mid X_2 = x_2$ for $x_2 \in \{-3, -2, -1, 0, 1, 2, 3\}$ (preferably on the same plot).

We will do this using R, first we will make a function for the bivariate PDF

```
# Joint PDF
f <- function(x1, x2){
  (1/(2*pi)) * (1 + x1^2 + x2^2)^(-3/2)
}
```

Recall that

$$f_{1|2}(x_1 \mid X_2 = x_2) = \frac{f(x_1, x_2)}{f_2(x_2)}.$$

So in order to plot the conditional distribution of $X_1 \mid X_2 = x_2$, we must first calculate $f_2(x_2)$. To do this we must integrate the joint PDF over x_1

$$\begin{aligned} f_2(x_2) &= \int_{-\infty}^{\infty} \frac{1}{2\pi}(1 + x_1^2 + x_2^2)^{-3/2} dx_1 \\ &= \frac{1}{2\pi} \int_{-\infty}^{\infty} (u)^{-3/2} dx_1 \\ &= \vdots \\ &= \frac{1}{\pi}(x_2^2 + 1)^{-1}. \end{aligned}$$

Similarly (this will be helpful when checking independence),

$$\begin{aligned} f_1(x_1) &= \int_{-\infty}^{\infty} \frac{1}{2\pi}(1 + x_2^2 + x_1^2)^{-3/2} dx_2 \\ &= \vdots \\ &= \frac{1}{\pi}(x_1^2 + 1)^{-1}. \end{aligned}$$

```


# Define the marginal PDFs f_1(x1) and f_2(x2) by integrating the joint PDF over
# x2 and x1 respectively
f_1 <- function(x1) {
  (1/pi) * (x1^2 + 1)^-1
}

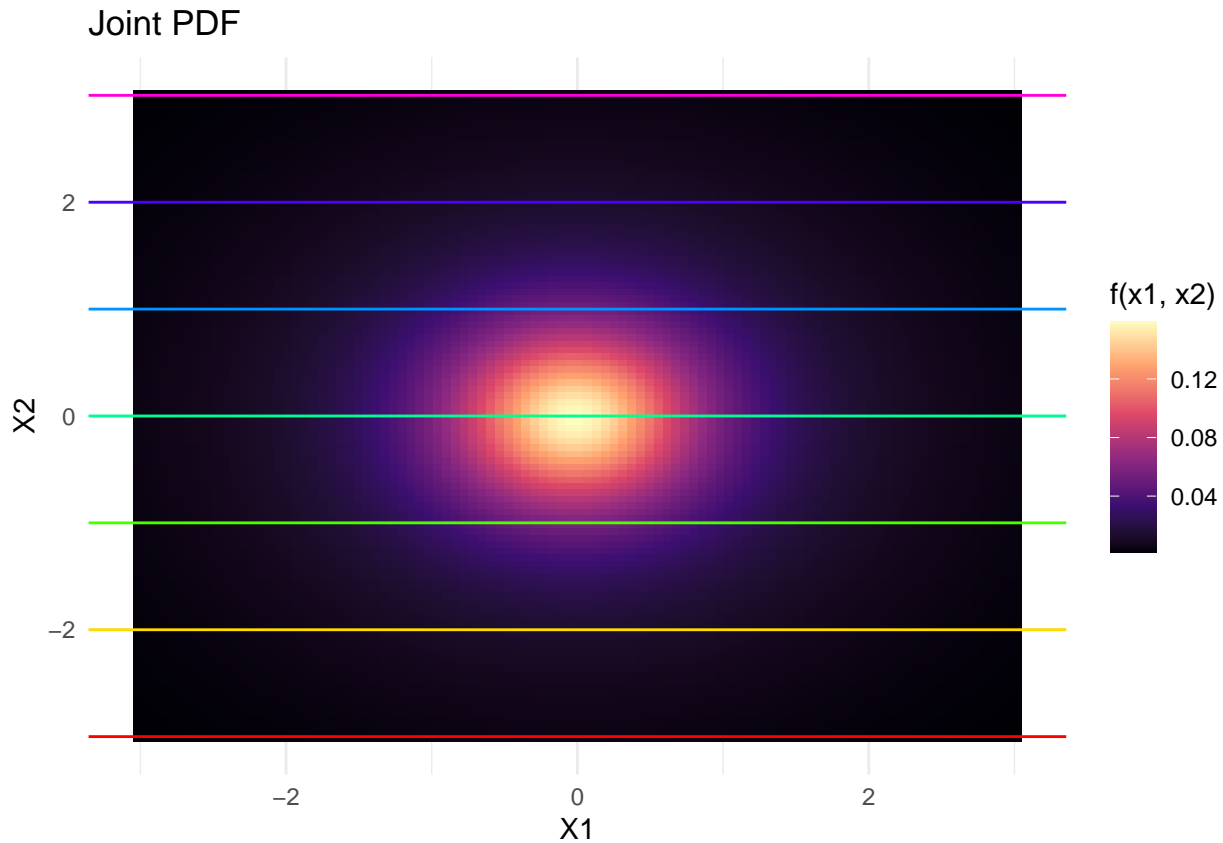
f_2 <- function(x2) {
  (1/pi) * (x2^2 + 1)^-1
}

# Define the conditional PDF of X1 given X2
f_conditional <- function(x1, x2){
  f(x1, x2)/f_2(x2)
}

```

Conditional Distributions of X1 given X2=x2





(b) Do X_1 and X_2 appear to be correlated? Justify your answer.

$$\mathbb{E}[X_1] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} x_1 f(x_1, x_2) dx_1 dx_2 = 0,$$

$$\mathbb{E}[X_2] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} x_2 f(x_1, x_2) dx_2 dx_1 = 0,$$

and

$$\mathbb{E}[X_1 X_2] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} x_1 x_2 f(x_1, x_2) dx_1 dx_2 = 0,$$

so

$$\text{Cov}[X_1, X_2] = \mathbb{E}[X_1 X_2] - \mathbb{E}[X_1] \mathbb{E}[X_2] = 0.$$

Since $\text{Cov}[X_1, X_2] = 0$, X_1 and X_2 are **uncorrelated**.

(c) Do X_1 and X_2 appear to be independent? Justify your answer.

```

marg_X1_vals <- sapply(x1_vals, f_1)
marg_X2_vals <- sapply(x2_vals, f_2)

# Check if f_1(x1) * f_2(x2) equals the joint PDF
check_independence <- sapply(1:length(x1_vals), function(i) {
  prod_vals <- f_1(x1_vals[i]) * f_2(x2_vals[i])
  joint_val <- f_conditional(x1_vals[i], x2_vals[i])
  return(abs(prod_vals - joint_val) < 1e-6) # tolerance for numerical error
})

all(check_independence) # TRUE if independent, FALSE otherwise

```

```
## [1] FALSE
```

The random variables X_1 and X_2 are **dependent** since the product of the marginal distributions does not equal the joint PDF. We can also see in the conditional distribution of $X_1 \mid X_2 = x_2$ plot that the spread of the bell curve **depends** on the value of x_2 .

(3) Chapter 1, Problem 9

For this problem pretend we are dealing with a language with a six-word dictionary

{fun, sun, sit, sat, fan, for}.

An extensive study of literature written in this language reveals that all words are equally likely except that “for” is α times as likely as the other words. Further study reveals that:

- i. Each keystroke is an error with probability θ .
- ii. All letters are equally likely to produce errors.
- iii. Given that a letter is typed incorrectly it is equally likely to be any other letter.
- iv. Errors are independent across letters.

For example, the probability of correctly typing “fun” (or any other word) is $(1 - \theta)^3$, the probability of typing “pun” or “fon” when intending to type is “fun” is $\theta(1 - \theta)^2$, and the probability of typing “foo” or “nnn” when intending to type “fun” is $\theta^2(1 - \theta)$. Use Bayes’ rule to develop a simple spell checker for this language.

```

spell_check <- function(typed_word, alpha, theta) {
  # Define the dictionary
  dictionary <- c("fun", "sun", "sit", "sat", "fan", "for")

  # Calculate the prior probabilities
  denom <- 5 + alpha
  Prob_for <- alpha / denom
  Prob_other <- 1 / denom

  # Define a function to calculate the likelihood P(t | w)
  # t is the typed word, w is the word in the dictionary

```

```

P_t_given_w <- function(t, w, theta) {
  matches <- sum(strsplit(t, NULL)[[1]] == strsplit(w, NULL)[[1]])
  k <- nchar(w)

  # Calculate likelihood based on number of matching letters
  if (matches == k) {
    return((1 - theta)^k)
  } else if (matches == k - 1) {
    return(theta * (1 - theta)^(k - 1))
  } else if (matches == k - 2) {
    return(theta^2 * (1 - theta)^(k - 2))
  } else if (matches == k - 3) {
    return(theta^3 * (1 - theta)^(k - 3))
  }
  return(0)
}

# Likelihoods for each word in the dictionary
likelihoods <- sapply(dictionary, function(w) P_t_given_w(typed_word, w, theta))

# Use Bayes Rule to calculate posterior probabilities
posteriors <- sapply(1:length(dictionary), function(i) {
  w <- dictionary[i]
  if (w == "for") {
    likelihoods[i] * Prob_for
  } else {
    likelihoods[i] * Prob_other
  }
})

# Normalize the posterior probabilities to sum to 1
posterior_sum <- sum(posteriors)
posteriors <- posteriors / posterior_sum

# Return the posterior probabilities
return(data.frame(posterior = posteriors))
}

```

For each of the typed words “sun”, “the”, “foo”, give the probability that each word in the dictionary was the intended word. Perform this for the parameters below:

(a) $\alpha = 2$ and $\theta = 0.1$.

```
spell_check("sun", alpha = 2, theta = 0.1)
```

```

##      posterior
## fun 0.09654350
## sun 0.86889154
## sit 0.01072706
## sat 0.01072706
## fan 0.01072706
## for 0.00238379

```

```
spell_check("the", alpha = 2, theta = 0.1)
```

```
##      posterior
## fun 0.1428571
## sun 0.1428571
## sit 0.1428571
## sat 0.1428571
## fan 0.1428571
## for 0.2857143
```

```
spell_check("foo", alpha = 2, theta = 0.1)
```

```
##      posterior
## fun 0.049180328
## sun 0.005464481
## sit 0.005464481
## sat 0.005464481
## fan 0.049180328
## for 0.885245902
```

(b) $\alpha = 50$ and $\theta = 0.1$.

```
spell_check("sun", alpha = 50, theta = 0.1)
```

```
##      posterior
## fun 0.09131905
## sun 0.82187148
## sit 0.01014656
## sat 0.01014656
## fan 0.01014656
## for 0.05636979
```

```
spell_check("the", alpha = 50, theta = 0.1)
```

```
##      posterior
## fun 0.01818182
## sun 0.01818182
## sit 0.01818182
## sat 0.01818182
## fan 0.01818182
## for 0.90909091
```

```
spell_check("foo", alpha = 50, theta = 0.1)
```

```
##      posterior
## fun 0.0022107590
## sun 0.0002456399
## sit 0.0002456399
## sat 0.0002456399
## fan 0.0022107590
## for 0.9948415623
```


(c) $\alpha = 2$ and $\theta = 0.95$.

```
spell_check("sun", alpha = 2, theta = 0.95)
```

```
##      posterior
## fun 1.281965e-03
## sun 6.747183e-05
## sit 2.435733e-02
## sat 2.435733e-02
## fan 2.435733e-02
## for 9.255786e-01
```

```
spell_check("the", alpha = 2, theta = 0.95)
```

```
##      posterior
## fun 0.1428571
## sun 0.1428571
## sit 0.1428571
## sat 0.1428571
## fan 0.1428571
## for 0.2857143
```

```
spell_check("foo", alpha = 2, theta = 0.95)
```

```
##      posterior
## fun 0.016918967
## sun 0.321460374
## sit 0.321460374
## sat 0.321460374
## fan 0.016918967
## for 0.001780944
```

Comment on the changes you observe in these three cases.

When the word “**sun**” is typed, which is in the dictionary, the spell checker correctly classifies it when $\theta = 0.1$. However, when $\theta = 0.95$, the spell checker thinks the user is typing “for”. This is because the keystroke error rate is set too high and “for” is α times as likely to occur.

By letter index, the word “**the**” has no matches with the words in the dictionary. In each case the spell checker classifies the typed word as “for” since it is α times as likely to occur.

The word “**foo**” is one letter off from “for.” The spell checker correctly classifies the word when $\theta = 0.1$. However, when $\theta = 0.95$, the keystroke error rate is set too high, so the spell checker thinks that typing an “f” is a mistake, as a result the spell checker misclassifies the word.

(4)

If 70% of a population is vaccinated, and the hospitalization rate is 5 times higher for an unvaccinated person than a vaccinated person, what is the probability that a person is vaccinated given they are hospitalized?

Given

$$\text{Prob}(\text{Vaccinated}) = 0.7 \implies \text{Prob}(\text{Unvaccinated}) = 0.3$$

and

$$\text{Prob}(\text{Hospitalized} \mid \text{Vaccinated}) = \frac{1}{5} \text{Prob}(\text{Hospitalized} \mid \text{Unvaccinated}).$$

Let

$$\text{Prob}(\text{Hospitalized} \mid \text{Vaccinated}) = r$$

and

$$\text{Prob}(\text{Hospitalized} \mid \text{Unvaccinated}) = 5r.$$

So

$$\text{Prob}(\text{Hospitalized}) = r \cdot 0.7 + 5r \cdot 0.3 = r \cdot 2.2.$$

We need to calculate $\text{Prob}(\text{Vaccinated} \mid \text{Hospitalized})$. This can be done using Bayes' rule.

$$\begin{aligned} \text{Prob}(\text{Vaccinated} \mid \text{Hospitalized}) &= \frac{\text{Prob}(\text{Vaccinated}) \text{Prob}(\text{Hospitalized} \mid \text{Vaccinated})}{\text{Prob}(\text{Hospitalized})} \\ &= \frac{0.7 \cdot r}{2.2 \cdot r} \\ &= \frac{0.7}{2.2} \\ &= 0.3182. \end{aligned}$$

Code Appendix

```
library(ggplot2)
library(viridis)
library(knitr)
# Joint PDF
f <- function(x1, x2){
  (1/(2*pi)) * (1 + x1^2 + x2^2)^(-3/2)
}
# Define the marginal PDFs f_1(x1) and f_2(x2) by integrating the joint PDF over
# x2 and x1 respectively
f_1 <- function(x1) {
  (1/pi) * (x1^2 + 1)^-1
}

f_2 <- function(x2) {
  (1/pi) * (x2^2 + 1)^-1
}

# Define the conditional PDF of X1 given X2
f_conditional <- function(x1, x2){
  f(x1, x2)/f_2(x2)
}
# Create a sequence of x1 values
x1_vals <- seq(-5, 5, length.out = 100)

# Define a vector of x2 values
x2_vals <- c(-3, -2, -1, 0, 1, 2, 3)

# Plot the conditional distributions for different x2 values
plot(NULL, xlim=c(-5, 5), ylim=c(0, .6), xlab="X1", ylab="Conditional PDF",
      main="Conditional Distributions of X1 given X2=x2")
colors <- rainbow(length(x2_vals))

for (i in 1:length(x2_vals)) {
  x2 <- x2_vals[i]
  y_vals <- sapply(x1_vals, function(x1) f_conditional(x1, x2))
  lines(x1_vals, y_vals, col=colors[i], lwd=2)
}

legend("topright", legend=paste("x2 =", x2_vals), col=colors, lwd=2)
# Set up a grid of values for x1 and x2
x1_vals <- seq(-3, 3, length.out = 100)
x2_vals <- seq(-3, 3, length.out = 100)

# Create a data frame for the joint PDF
joint_pdf_vals <- expand.grid(x1 = x1_vals, x2 = x2_vals)
joint_pdf_vals$z <- mapply(f, joint_pdf_vals$x1, joint_pdf_vals$x2)
# Plot the joint PDF as a contour plot
ggplot(joint_pdf_vals, aes(x = x1, y = x2, z = z)) +
  geom_tile(aes(fill = z), width = 0.1, height = 0.1) +
  scale_fill_viridis_c(option = "magma", name = "f(x1, x2)") +
  ggtitle("Joint PDF") +
```

```

labs(x = "X1", y = "X2") +
geom_hline(yintercept = -3, color = colors[1]) +
geom_hline(yintercept = -2, color = colors[2]) +
geom_hline(yintercept = -1, color = colors[3]) +
geom_hline(yintercept = 0, color = colors[4]) +
geom_hline(yintercept = 1, color = colors[5]) +
geom_hline(yintercept = 2, color = colors[6]) +
geom_hline(yintercept = 3, color = colors[7]) +
theme_minimal()
marg_X1_vals <- sapply(x1_vals, f_1)
marg_X2_vals <- sapply(x2_vals, f_2)

# Check if f_1(x1) * f_2(x2) equals the joint PDF
check_independence <- sapply(1:length(x1_vals), function(i) {
  prod_vals <- f_1(x1_vals[i]) * f_2(x2_vals[i])
  joint_val <- f_conditional(x1_vals[i], x2_vals[i])
  return(abs(prod_vals - joint_val) < 1e-6) # tolerance for numerical error
})

all(check_independence) # TRUE if independent, FALSE otherwise
spell_check <- function(typed_word, alpha, theta) {
  # Define the dictionary
  dictionary <- c("fun", "sun", "sit", "sat", "fan", "for")

  # Calculate the prior probabilities
  denom <- 5 + alpha
  Prob_for <- alpha / denom
  Prob_other <- 1 / denom

  # Define a function to calculate the likelihood P(t | w)
  # t is the typed word, w is the word in the dictionary
  P_t_given_w <- function(t, w, theta) {
    matches <- sum(strsplit(t, NULL)[[1]] == strsplit(w, NULL)[[1]])
    k <- nchar(w)

    # Calculate likelihood based on number of matching letters
    if (matches == k) {
      return((1 - theta)^k)
    } else if (matches == k - 1) {
      return(theta * (1 - theta)^(k - 1))
    } else if (matches == k - 2) {
      return(theta^2 * (1 - theta)^(k - 2))
    } else if (matches == k - 3) {
      return(theta^3 * (1 - theta)^(k - 3))
    }
    return(0)
  }

  # Likelihoods for each word in the dictionary
  likelihoods <- sapply(dictionary, function(w) P_t_given_w(typed_word, w, theta))

  # Use Bayes Rule to calculate posterior probabilities
  posteriors <- sapply(1:length(dictionary), function(i) {

```

```

    w <- dictionary[i]
    if (w == "for") {
      likelihoods[i] * Prob_for
    } else {
      likelihoods[i] * Prob_other
    }
  })

  # Normalize the posterior probabilities to sum to 1
  posterior_sum <- sum(posterior)
  posterior <- posterior / posterior_sum

  # Return the posterior probabilities
  return(data.frame(posterior = posterior))
}

spell_check("sun", alpha = 2, theta = 0.1)
spell_check("the", alpha = 2, theta = 0.1)
spell_check("foo", alpha = 2, theta = 0.1)
spell_check("sun", alpha = 50, theta = 0.1)
spell_check("the", alpha = 50, theta = 0.1)
spell_check("foo", alpha = 50, theta = 0.1)
spell_check("sun", alpha = 2, theta = 0.95)
spell_check("the", alpha = 2, theta = 0.95)
spell_check("foo", alpha = 2, theta = 0.95)

```